

# finding\_donors

June 15, 2021

## 0.1 Supervised Learning

## 0.2 Project: Finding Donors for *CharityML*

In this notebook, some template code has already been provided for you, and it will be your job to implement the additional functionality necessary to successfully complete this project. Sections that begin with '**Implementation**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section and the specifics of the implementation are marked in the code block with a '`TODO`' statement. Please be sure to read the instructions carefully!

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Please specify WHICH VERSION OF PYTHON you are using when submitting this notebook. Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

## 0.3 Getting Started

In this project, you will employ several supervised algorithms of your choice to accurately model individuals' income using data collected from the 1994 U.S. Census. You will then choose the best candidate algorithm from preliminary results and further optimize this algorithm to best model the data. Your goal with this implementation is to construct a model that accurately predicts whether an individual makes more than \$50,000. This sort of task can arise in a non-profit setting, where organizations survive on donations. Understanding an individual's income can help a non-profit better understand how large of a donation to request, or whether or not they should reach out to begin with. While it can be difficult to determine an individual's general income bracket directly from public sources, we can (as we will see) infer this value from other publically available features.

The dataset for this project originates from the [UCI Machine Learning Repository](#). The dataset was donated by Ron Kohavi and Barry Becker, after being published in the article "*Scaling Up the Accuracy of Naive-Bayes Classifiers: A Decision-Tree Hybrid*". You can find the article by Ron Kohavi [online](#). The data we investigate here consists of small changes to the original dataset, such as removing the '`fnlwgt`' feature and records with missing or ill-formatted entries.

---

## 0.4 Exploring the Data

Run the code cell below to load necessary Python libraries and load the census data. Note that the last column from this dataset, 'income', will be our target label (whether an individual makes more than, or at most, \$50,000 annually). All other columns are features about each individual in the census database.

```
In [27]: # Import libraries necessary for this project
import numpy as np
import pandas as pd
from time import time
from IPython.display import display # Allows the use of display() for DataFrames

# Import supplementary visualization code visuals.py
import visuals as vs

# Pretty display for notebooks
%matplotlib inline

# Load the Census dataset
data = pd.read_csv("census.csv")

# Success - Display the first record
display(data.head(n=1))
```

	age	workclass	education_level	education-num	marital-status	\
0	39	State-gov	Bachelors	13.0	Never-married	
	occupation	relationship	race	sex	capital-gain	capital-loss \
0	Adm-clerical	Not-in-family	White	Male	2174.0	0.0
	hours-per-week	native-country	income			
0	40.0	United-States	<=50K			

### 0.4.1 Implementation: Data Exploration

A cursory investigation of the dataset will determine how many individuals fit into either group, and will tell us about the percentage of these individuals making more than \$50,000. In the code cell below, you will need to compute the following: - The total number of records, 'n\_records' - The number of individuals making more than \$50,000 annually, 'n\_greater\_50k'. - The number of individuals making at most \$50,000 annually, 'n\_at\_most\_50k'. - The percentage of individuals making more than \$50,000 annually, 'greater\_percent'.

**\*\* HINT: \*\*** You may need to look at the table above to understand how the 'income' entries are formatted.

```
In [28]: data
```

```

Out[28]:
   age  workclass  education_level  education-num  \
0    39    State-gov    Bachelors          13.0
1    50  Self-emp-not-inc    Bachelors          13.0
2    38    Private    HS-grad           9.0
3    53    Private    11th              7.0
4    28    Private    Bachelors          13.0
5    37    Private    Masters          14.0
6    49    Private    9th              5.0
7    52  Self-emp-not-inc    HS-grad           9.0
8    31    Private    Masters          14.0
9    42    Private    Bachelors          13.0
10   37    Private    Some-college        10.0
11   30    State-gov    Bachelors          13.0
12   23    Private    Bachelors          13.0
13   32    Private    Assoc-acdm         12.0
14   34    Private    7th-8th           4.0
15   25  Self-emp-not-inc    HS-grad           9.0
16   32    Private    HS-grad           9.0
17   38    Private    11th              7.0
18   43  Self-emp-not-inc    Masters          14.0
19   40    Private    Doctorate          16.0
20   54    Private    HS-grad           9.0
21   35    Federal-gov    9th            5.0
22   43    Private    11th              7.0
23   59    Private    HS-grad           9.0
24   56    Local-gov    Bachelors          13.0
25   19    Private    HS-grad           9.0
26   39    Private    HS-grad           9.0
27   49    Private    HS-grad           9.0
28   23    Local-gov    Assoc-acdm         12.0
29   20    Private    Some-college        10.0
...   ...   ...   ...   ...
45192  25    Private    HS-grad           9.0
45193  31    Private    HS-grad           9.0
45194  49    Self-emp-inc    HS-grad           9.0
45195  60    Private    Assoc-voc         11.0
45196  39    Private    Bachelors          13.0
45197  38    Private    Masters          14.0
45198  43    Local-gov    Masters          14.0
45199  23    Private    HS-grad           9.0
45200  73    Self-emp-inc    Some-college        10.0
45201  35    Private    Some-college        10.0
45202  66    Private    HS-grad           9.0
45203  27    Private    Some-college        10.0
45204  40    Private    Prof-school         15.0
45205  51    Private    HS-grad           9.0
45206  22    Private    Some-college        10.0
45207  64  Self-emp-not-inc    HS-grad           9.0

```

45208	55	Private	HS-grad	9.0
45209	38	Private	Assoc-voc	11.0
45210	58	Private	Assoc-acdm	12.0
45211	32	Private	HS-grad	9.0
45212	48	Private	HS-grad	9.0
45213	61	Private	HS-grad	9.0
45214	31	Private	HS-grad	9.0
45215	25	Private	HS-grad	9.0
45216	48	Local-gov	Masters	14.0
45217	33	Private	Bachelors	13.0
45218	39	Private	Bachelors	13.0
45219	38	Private	Bachelors	13.0
45220	44	Private	Bachelors	13.0
45221	35	Self-emp-inc	Bachelors	13.0

	marital-status	occupation	relationship \
0	Never-married	Adm-clerical	Not-in-family
1	Married-civ-spouse	Exec-managerial	Husband
2	Divorced	Handlers-cleaners	Not-in-family
3	Married-civ-spouse	Handlers-cleaners	Husband
4	Married-civ-spouse	Prof-specialty	Wife
5	Married-civ-spouse	Exec-managerial	Wife
6	Married-spouse-absent	Other-service	Not-in-family
7	Married-civ-spouse	Exec-managerial	Husband
8	Never-married	Prof-specialty	Not-in-family
9	Married-civ-spouse	Exec-managerial	Husband
10	Married-civ-spouse	Exec-managerial	Husband
11	Married-civ-spouse	Prof-specialty	Husband
12	Never-married	Adm-clerical	Own-child
13	Never-married	Sales	Not-in-family
14	Married-civ-spouse	Transport-moving	Husband
15	Never-married	Farming-fishing	Own-child
16	Never-married	Machine-op-inspct	Unmarried
17	Married-civ-spouse	Sales	Husband
18	Divorced	Exec-managerial	Unmarried
19	Married-civ-spouse	Prof-specialty	Husband
20	Separated	Other-service	Unmarried
21	Married-civ-spouse	Farming-fishing	Husband
22	Married-civ-spouse	Transport-moving	Husband
23	Divorced	Tech-support	Unmarried
24	Married-civ-spouse	Tech-support	Husband
25	Never-married	Craft-repair	Own-child
26	Divorced	Exec-managerial	Not-in-family
27	Married-civ-spouse	Craft-repair	Husband
28	Never-married	Protective-serv	Not-in-family
29	Never-married	Sales	Own-child
...	...	...	...
45192	Divorced	Machine-op-inspct	Not-in-family

45193	Never-married	Machine-op-inspct	Not-in-family
45194	Married-civ-spouse	Exec-managerial	Husband
45195	Married-civ-spouse	Prof-specialty	Husband
45196	Never-married	Tech-support	Not-in-family
45197	Married-civ-spouse	Prof-specialty	Husband
45198	Married-civ-spouse	Exec-managerial	Husband
45199	Never-married	Machine-op-inspct	Own-child
45200	Divorced	Exec-managerial	Not-in-family
45201	Married-civ-spouse	Protective-serv	Husband
45202	Widowed	Sales	Other-relative
45203	Never-married	Sales	Not-in-family
45204	Married-civ-spouse	Prof-specialty	Husband
45205	Married-civ-spouse	Craft-repair	Husband
45206	Never-married	Craft-repair	Own-child
45207	Widowed	Farming-fishing	Not-in-family
45208	Separated	Priv-house-serv	Not-in-family
45209	Never-married	Adm-clerical	Unmarried
45210	Divorced	Prof-specialty	Not-in-family
45211	Married-civ-spouse	Handlers-cleaners	Husband
45212	Married-civ-spouse	Adm-clerical	Husband
45213	Married-civ-spouse	Sales	Husband
45214	Married-civ-spouse	Craft-repair	Husband
45215	Never-married	Other-service	Own-child
45216	Divorced	Other-service	Not-in-family
45217	Never-married	Prof-specialty	Own-child
45218	Divorced	Prof-specialty	Not-in-family
45219	Married-civ-spouse	Prof-specialty	Husband
45220	Divorced	Adm-clerical	Own-child
45221	Married-civ-spouse	Exec-managerial	Husband

	race	sex	capital-gain	capital-loss	\
0	White	Male	2174.0	0.0	
1	White	Male	0.0	0.0	
2	White	Male	0.0	0.0	
3	Black	Male	0.0	0.0	
4	Black	Female	0.0	0.0	
5	White	Female	0.0	0.0	
6	Black	Female	0.0	0.0	
7	White	Male	0.0	0.0	
8	White	Female	14084.0	0.0	
9	White	Male	5178.0	0.0	
10	Black	Male	0.0	0.0	
11	Asian-Pac-Islander	Male	0.0	0.0	
12	White	Female	0.0	0.0	
13	Black	Male	0.0	0.0	
14	Amer-Indian-Eskimo	Male	0.0	0.0	
15	White	Male	0.0	0.0	
16	White	Male	0.0	0.0	

17	White	Male	0.0	0.0
18	White	Female	0.0	0.0
19	White	Male	0.0	0.0
20	Black	Female	0.0	0.0
21	Black	Male	0.0	0.0
22	White	Male	0.0	2042.0
23	White	Female	0.0	0.0
24	White	Male	0.0	0.0
25	White	Male	0.0	0.0
26	White	Male	0.0	0.0
27	White	Male	0.0	0.0
28	White	Male	0.0	0.0
29	Black	Male	0.0	0.0
...	...	...	...	...
45192	Black	Male	0.0	0.0
45193	White	Male	0.0	0.0
45194	White	Male	0.0	0.0
45195	White	Male	7688.0	0.0
45196	White	Female	0.0	1669.0
45197	White	Male	0.0	0.0
45198	White	Male	0.0	1902.0
45199	White	Male	0.0	0.0
45200	White	Female	0.0	0.0
45201	White	Male	0.0	0.0
45202	White	Female	0.0	0.0
45203	White	Female	0.0	0.0
45204	White	Male	15024.0	0.0
45205	White	Male	0.0	0.0
45206	White	Male	0.0	0.0
45207	White	Male	0.0	0.0
45208	White	Female	0.0	0.0
45209	Black	Female	0.0	0.0
45210	White	Male	0.0	0.0
45211	White	Male	0.0	0.0
45212	White	Male	0.0	0.0
45213	White	Male	0.0	0.0
45214	White	Male	0.0	0.0
45215	White	Female	0.0	0.0
45216	White	Male	0.0	0.0
45217	White	Male	0.0	0.0
45218	White	Female	0.0	0.0
45219	White	Male	0.0	0.0
45220	Asian-Pac-Islander	Male	5455.0	0.0
45221	White	Male	0.0	0.0

	hours-per-week	native-country	income
0	40.0	United-States	<=50K
1	13.0	United-States	<=50K

2	40.0	United-States	<=50K
3	40.0	United-States	<=50K
4	40.0	Cuba	<=50K
5	40.0	United-States	<=50K
6	16.0	Jamaica	<=50K
7	45.0	United-States	>50K
8	50.0	United-States	>50K
9	40.0	United-States	>50K
10	80.0	United-States	>50K
11	40.0	India	>50K
12	30.0	United-States	<=50K
13	50.0	United-States	<=50K
14	45.0	Mexico	<=50K
15	35.0	United-States	<=50K
16	40.0	United-States	<=50K
17	50.0	United-States	<=50K
18	45.0	United-States	>50K
19	60.0	United-States	>50K
20	20.0	United-States	<=50K
21	40.0	United-States	<=50K
22	40.0	United-States	<=50K
23	40.0	United-States	<=50K
24	40.0	United-States	>50K
25	40.0	United-States	<=50K
26	80.0	United-States	<=50K
27	40.0	United-States	<=50K
28	52.0	United-States	<=50K
29	44.0	United-States	<=50K
...	...	...	...
45192	40.0	United-States	<=50K
45193	40.0	United-States	<=50K
45194	40.0	Canada	>50K
45195	40.0	United-States	>50K
45196	40.0	United-States	<=50K
45197	50.0	United-States	>50K
45198	50.0	United-States	>50K
45199	40.0	United-States	<=50K
45200	40.0	United-States	<=50K
45201	40.0	United-States	<=50K
45202	8.0	United-States	<=50K
45203	45.0	United-States	<=50K
45204	55.0	United-States	>50K
45205	40.0	United-States	<=50K
45206	40.0	United-States	<=50K
45207	32.0	United-States	<=50K
45208	32.0	United-States	<=50K
45209	40.0	United-States	<=50K
45210	36.0	United-States	<=50K

45211	40.0	United-States	<=50K
45212	40.0	United-States	<=50K
45213	48.0	United-States	<=50K
45214	40.0	United-States	<=50K
45215	40.0	United-States	<=50K
45216	40.0	United-States	<=50K
45217	40.0	United-States	<=50K
45218	36.0	United-States	<=50K
45219	50.0	United-States	<=50K
45220	40.0	United-States	<=50K
45221	60.0	United-States	>50K

[45222 rows x 14 columns]

In [29]: data.dtypes

```
Out[29]: age                int64
workclass                object
education_level          object
education-num            float64
marital-status           object
occupation               object
relationship             object
race                     object
sex                      object
capital-gain             float64
capital-loss             float64
hours-per-week           float64
native-country           object
income                   object
dtype: object
```

In [30]: *# TODO: Total number of records*

```
n_records = data.shape[0]
```

*# TODO: Number of records where individual's income is more than \$50,000*

```
more = data[data['income'] == '>50K']
```

```
n_greater_50k = more.shape[0]
```

*# TODO: Number of records where individual's income is at most \$50,000*

```
less = data[data['income'] == '<=50K']
```

```
n_at_most_50k = less.shape[0]
```

*# TODO: Percentage of individuals whose income is more than \$50,000*

```
greater_percent = (n_greater_50k/n_records)*100
```

*# Print the results*

```
print("Total number of records: {}".format(n_records))
```



```

print("Individuals making more than $50,000: {}".format(n_greater_50k))
print("Individuals making at most $50,000: {}".format(n_at_most_50k))
print("Percentage of individuals making more than $50,000: {}%".format(greater_percent))

```

Total number of records: 45222

Individuals making more than \$50,000: 11208

Individuals making at most \$50,000: 34014

Percentage of individuals making more than \$50,000: 24.78439697492371%

## **\*\* Featureset Exploration \*\***

- **age**: continuous.
- **workclass**: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- **education**: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- **education-num**: continuous.
- **marital-status**: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- **occupation**: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- **relationship**: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- **race**: Black, White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other.
- **sex**: Female, Male.
- **capital-gain**: continuous.
- **capital-loss**: continuous.
- **hours-per-week**: continuous.
- **native-country**: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

---

## **0.5 Preparing the Data**

Before data can be used as input for machine learning algorithms, it often must be cleaned, formatted, and restructured — this is typically known as **preprocessing**. Fortunately, for this dataset, there are no invalid or missing entries we must deal with, however, there are some qualities about certain features that must be adjusted. This preprocessing can help tremendously with the outcome and predictive power of nearly all learning algorithms.

### **0.5.1 Transforming Skewed Continuous Features**


A dataset may sometimes contain at least one feature whose values tend to lie near a single number, but will also have a non-trivial number of vastly larger or smaller values than that single

number. Algorithms can be sensitive to such distributions of values and can underperform if the range is not properly normalized. With the census dataset two features fit this description: 'capital-gain' and 'capital-loss'.

Run the code cell below to plot a histogram of these two features. Note the range of the values present and how they are distributed.

```
In [31]: # Split the data into features and target label
income_raw = data['income']
features_raw = data.drop('income', axis = 1)

# Visualize skewed continuous features of original data
vs.distribution(data)
```




For highly-skewed feature distributions such as 'capital-gain' and 'capital-loss', it is common practice to apply a logarithmic transformation on the data so that the very large and very small values do not negatively affect the performance of a learning algorithm. Using a logarithmic transformation significantly reduces the range of values caused by outliers. Care must be taken when applying this transformation however: The logarithm of 0 is undefined, so we must translate the values by a small amount above 0 to apply the the logarithm successfully.

Run the code cell below to perform a transformation on the data and visualize the results. Again, note the range of values and how they are distributed.

```
In [32]: # Log-transform the skewed features
skewed = ['capital-gain', 'capital-loss']
features_log_transformed = pd.DataFrame(data = features_raw)
features_log_transformed[skewed] = features_raw[skewed].apply(lambda x: np.log(x + 1))

# Visualize the new log distributions
vs.distribution(features_log_transformed, transformed = True)
```



## 0.5.2 Normalizing Numerical Features

In addition to performing transformations on features that are highly skewed, it is often good practice to perform some type of scaling on numerical features. Applying a scaling to the data does not change the shape of each feature's distribution (such as 'capital-gain' or 'capital-loss' above); however, normalization ensures that each feature is treated equally when applying supervised learners. Note that once scaling is applied, observing the data in its raw form will no longer have the same original meaning, as exemplified below.

Run the code cell below to normalize each numerical feature. We will use `sklearn.preprocessing.MinMaxScaler` for this.

```
In [33]: # Import sklearn.preprocessing.StandardScaler
         from sklearn.preprocessing import MinMaxScaler

         # Initialize a scaler, then apply it to the features
         scaler = MinMaxScaler() # default=(0, 1)
         numerical = ['age', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']

         features_log_minmax_transform = pd.DataFrame(data = features_log_transformed)
         features_log_minmax_transform[numerical] = scaler.fit_transform(features_log_transformed[numerical])
         # Show an example of a record with scaling applied
         display(features_log_minmax_transform.head(n = 20))
```

	age	workclass	education_level	education-num \
0	0.301370	State-gov	Bachelors	0.800000
1	0.452055	Self-emp-not-inc	Bachelors	0.800000
2	0.287671	Private	HS-grad	0.533333
3	0.493151	Private	11th	0.400000
4	0.150685	Private	Bachelors	0.800000
5	0.273973	Private	Masters	0.866667

6	0.438356	Private	9th	0.266667
7	0.479452	Self-emp-not-inc	HS-grad	0.533333
8	0.191781	Private	Masters	0.866667
9	0.342466	Private	Bachelors	0.800000
10	0.273973	Private	Some-college	0.600000
11	0.178082	State-gov	Bachelors	0.800000
12	0.082192	Private	Bachelors	0.800000
13	0.205479	Private	Assoc-acdm	0.733333
14	0.232877	Private	7th-8th	0.200000
15	0.109589	Self-emp-not-inc	HS-grad	0.533333
16	0.205479	Private	HS-grad	0.533333
17	0.287671	Private	11th	0.400000
18	0.356164	Self-emp-not-inc	Masters	0.866667
19	0.315068	Private	Doctorate	1.000000

	marital-status	occupation	relationship \
0	Never-married	Adm-clerical	Not-in-family
1	Married-civ-spouse	Exec-managerial	Husband
2	Divorced	Handlers-cleaners	Not-in-family
3	Married-civ-spouse	Handlers-cleaners	Husband
4	Married-civ-spouse	Prof-specialty	Wife
5	Married-civ-spouse	Exec-managerial	Wife
6	Married-spouse-absent	Other-service	Not-in-family
7	Married-civ-spouse	Exec-managerial	Husband
8	Never-married	Prof-specialty	Not-in-family
9	Married-civ-spouse	Exec-managerial	Husband
10	Married-civ-spouse	Exec-managerial	Husband
11	Married-civ-spouse	Prof-specialty	Husband
12	Never-married	Adm-clerical	Own-child
13	Never-married	Sales	Not-in-family
14	Married-civ-spouse	Transport-moving	Husband
15	Never-married	Farming-fishing	Own-child
16	Never-married	Machine-op-inspct	Unmarried
17	Married-civ-spouse	Sales	Husband
18	Divorced	Exec-managerial	Unmarried
19	Married-civ-spouse	Prof-specialty	Husband

	race	sex	capital-gain	capital-loss	hours-per-week \
0	White	Male	0.667492	0.0	0.397959
1	White	Male	0.000000	0.0	0.122449
2	White	Male	0.000000	0.0	0.397959
3	Black	Male	0.000000	0.0	0.397959
4	Black	Female	0.000000	0.0	0.397959
5	White	Female	0.000000	0.0	0.397959
6	Black	Female	0.000000	0.0	0.153061
7	White	Male	0.000000	0.0	0.448980
8	White	Female	0.829751	0.0	0.500000
9	White	Male	0.742849	0.0	0.397959

10		Black	Male	0.000000	0.0	0.806122
11	Asian-Pac-Islander		Male	0.000000	0.0	0.397959
12		White	Female	0.000000	0.0	0.295918
13		Black	Male	0.000000	0.0	0.500000
14	Amer-Indian-Eskimo		Male	0.000000	0.0	0.448980
15		White	Male	0.000000	0.0	0.346939
16		White	Male	0.000000	0.0	0.397959
17		White	Male	0.000000	0.0	0.500000
18		White	Female	0.000000	0.0	0.448980
19		White	Male	0.000000	0.0	0.602041

	native-country
0	United-States
1	United-States
2	United-States
3	United-States
4	Cuba
5	United-States
6	Jamaica
7	United-States
8	United-States
9	United-States
10	United-States
11	India
12	United-States
13	United-States
14	Mexico
15	United-States
16	United-States
17	United-States
18	United-States
19	United-States

### 0.5.3 Implementation: Data Preprocessing

From the table in **Exploring the Data** above, we can see there are several features for each record that are non-numeric. Typically, learning algorithms expect input to be numeric, which requires that non-numeric features (called *categorical variables*) be converted. One popular way to convert categorical variables is by using the **one-hot encoding** scheme. One-hot encoding creates a “dummy” variable for each possible category of each non-numeric feature. For example, assume someFeature has three possible entries: A, B, or C. We then encode this feature into someFeature\_A, someFeature\_B and someFeature\_C.

	someFeature	someFeature_A	someFeature_B	someFeature_C
0	B	0	1	0
1	C	0	0	1
2	A	1	0	0

one-hot encode

Additionally, as with the non-numeric features, we need to convert the non-numeric target label, 'income' to numerical values for the learning algorithm to work. Since there are only two possible categories for this label (" $\leq 50K$ " and " $> 50K$ "), we can avoid using one-hot encoding and simply encode these two categories as 0 and 1, respectively. In code cell below, you will need to implement the following:

- Use `pandas.get_dummies()` to perform one-hot encoding on the 'features\_log\_minmax\_transform' data.
- Convert the target label 'income\_raw' to numerical entries.
- Set records with " $\leq 50K$ " to 0 and records with " $> 50K$ " to 1.

```
In [34]: # TODO: One-hot encode the 'features_log_minmax_transform' data using pandas.get_dummies  
features_final = pd.get_dummies(features_log_minmax_transform)  
  
# TODO: Encode the 'income_raw' data to numerical values  
income = income_raw.map({"<=50K": 0, ">50K": 1})  
  
# Print the number of features after one-hot encoding  
encoded = list(features_final.columns)  
print("{} total features after one-hot encoding.".format(len(encoded)))  
  
# Uncomment the following line to see the encoded feature names  
print(encoded)
```

```
103 total features after one-hot encoding.  
['age', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week', 'workclass_Federal-gov']
```

### 0.5.4 Shuffle and Split Data

Now all *categorical variables* have been converted into numerical features, and all numerical features have been normalized. As always, we will now split the data (both features and their labels) into training and test sets. 80% of the data will be used for training and 20% for testing.

Run the code cell below to perform this split.

```
In [35]: # Import train_test_split
         from sklearn.cross_validation import train_test_split

         # Split the 'features' and 'income' data into training and testing sets
         X_train, X_test, y_train, y_test = train_test_split(features_final,
                                                             income,
                                                             test_size = 0.2,
                                                             random_state = 0)

         # Show the results of the split
         print("Training set has {} samples.".format(X_train.shape[0]))
         print("Testing set has {} samples.".format(X_test.shape[0]))
```

```
Training set has 36177 samples.
Testing set has 9045 samples.
```

*Note: this Workspace is running on sklearn v0.19. If you use the newer version ( $\geq 0.20$ ), the `sklearn.cross_validation` has been replaced with `sklearn.model_selection`.*

---

## 0.6 Evaluating Model Performance

In this section, we will investigate four different algorithms, and determine which is best at modeling the data. Three of these algorithms will be supervised learners of your choice, and the fourth algorithm is known as a *naive predictor*.

### 0.6.1 Metrics and the Naive Predictor

*CharityML*, equipped with their research, knows individuals that make more than \$50,000 are most likely to donate to their charity. Because of this, *CharityML* is particularly interested in predicting who makes more than \$50,000 accurately. It would seem that using **accuracy** as a metric for evaluating a particular model's performance would be appropriate. Additionally, identifying someone that *does not* make more than \$50,000 as someone who does would be detrimental to *CharityML*, since they are looking to find individuals willing to donate. Therefore, a model's ability to precisely predict those that make more than \$50,000 is *more important* than the model's ability to **recall** those individuals. We can use **F-beta score** as a metric that considers both precision and recall:

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

In particular, when  $\beta = 0.5$ , more emphasis is placed on precision. This is called the **F<sub>0.5</sub> score** (or F-score for simplicity).

Looking at the distribution of classes (those who make at most \$50,000, and those who make more), it's clear most individuals do not make more than \$50,000. This can greatly affect **accuracy**, since we could simply say "*this person does not make more than \$50,000*" and generally be right, without ever looking at the data! Making such a statement would be called **naive**, since we have not considered any information to substantiate the claim. It is always important to consider the *naive prediction* for your data, to help establish a benchmark for whether a model is performing well. That been said, using that prediction would be pointless: If we predicted all people made less than \$50,000, *CharityML* would identify no one as donors.

**Note: Recap of accuracy, precision, recall** \*\* Accuracy \*\* measures how often the classifier makes the correct prediction. It's the ratio of the number of correct predictions to the total number of predictions (the number of test data points).

\*\* Precision \*\* tells us what proportion of messages we classified as spam, actually were spam. It is a ratio of true positives(words classified as spam, and which are actually spam) to all positives(all words classified as spam, irrespective of whether that was the correct classificatio), in other words it is the ratio of

[True Positives/(True Positives + False Positives)]

\*\* Recall(sensitivity)\*\* tells us what proportion of messages that actually were spam were classified by us as spam. It is a ratio of true positives(words classified as spam, and which are actually spam) to all the words that were actually spam, in other words it is the ratio of

[True Positives/(True Positives + False Negatives)]

For classification problems that are skewed in their classification distributions like in our case, for example if we had a 100 text messages and only 2 were spam and the rest 98 weren't, accuracy by itself is not a very good metric. We could classify 90 messages as not spam(including the 2 that were spam but we classify them as not spam, hence they would be false negatives) and 10 as spam(all 10 false positives) and still get a reasonably good accuracy score. For such cases, precision and recall come in very handy. These two metrics can be combined to get the F1 score, which is weighted average(harmonic mean) of the precision and recall scores. This score can range from 0 to 1, with 1 being the best possible F1 score(we take the harmonic mean as we are dealing with ratios).

### 0.6.2 Question 1 - Naive Predictor Performace

- If we chose a model that always predicted an individual made more than \$50,000, what would that model's accuracy and F-score be on this dataset? You must use the code cell below and assign your results to 'accuracy' and 'fscore' to be used later.

**\*\* Please note \*\*** that the the purpose of generating a naive predictor is simply to show what a base model without any intelligence would look like. In the real world, ideally your base model would be either the results of a previous model or could be based on a research paper upon which you are looking to improve. When there is no benchmark model set, getting a result better than random choice is a place you could start from.

**\*\* HINT: \*\***

- When we have a model that always predicts '1' (i.e. the individual makes more than 50k) then our model will have no True Negatives(TN) or False Negatives(FN) as we are not making any negative('0' value) predictions. Therefore our Accuracy in this case becomes the same as our Precision(True Positives/(True Positives + False Positives)) as every prediction that we have made with value '1' that should have '0' becomes a False Positive; therefore our denominator in this case is the total number of records we have in total.
- Our Recall score(True Positives/(True Positives + False Negatives)) in this setting becomes 1 as we have no False Negatives.

```
In [36]: TP = np.sum(income) # Counting the ones as this is the naive case. Note that 'income' is
#encoded to numerical values done in the data preprocessing step.
FP = income.count() - TP # Specific to the naive case

TN = 0 # No predicted negatives in the naive case
FN = 0 # No predicted negatives in the naive case

# TODO: Calculate accuracy, precision and recall
accuracy = (TP+TN)/(TP+TN+FP+FN)
recall = (TP)/(TP+FN)
precision = (TP)/(TP+FP)

# TODO: Calculate F-score using the formula above for beta = 0.5 and correct values for
fscore = ((1+((0.5)**2))*(precision*recall))/((precision*((0.5)**2))+recall)

# Print the results
print("Naive Predictor: [Accuracy score: {:.4f}, F-score: {:.4f}]"
```



Naive Predictor: [Accuracy score: 0.2478, F-score: 0.2917]

### 0.6.3 Supervised Learning Models

The following are some of the supervised learning models that are currently available in [scikit-learn](#) that you may choose from: - Gaussian Naive Bayes (GaussianNB) - Decision Trees - Ensemble Methods (Bagging, AdaBoost, Random Forest, Gradient Boosting) - K-Nearest Neighbors (KNeighbors) - Stochastic Gradient Descent Classifier (SGDC) - Support Vector Machines (SVM) - Logistic Regression

### 0.6.4 Question 2 - Model Application

List three of the supervised learning models above that are appropriate for this problem that you will test on the census data. For each model chosen

- Describe one real-world application in industry where the model can be applied.
- What are the strengths of the model; when does it perform well?
- What are the weaknesses of the model; when does it perform poorly?
- What makes this model a good candidate for the problem, given what you know about the data?

**HINT:**

Structure your answer in the same format as above, with 4 parts for each of the three models you pick. Please include references with your answer.

**Answer:**

1-Logistic Regression

- we can apply Logistic Regression to classify plants according to their features
- the model will work well in the data which has lots of labels
- the model won't work well in a nonlinear relationship between independent and dependant variables

Overall, Logistic Regression is a good choice since the labels are binary and the problem seems linear.

2- Decision Tree

- this model can use to classify electronic devices according to there shape, size, trade mark...etc
- the model is easy to implement and understand even for non-technical people
- the model could not work well if the data have a tiny variety

Overall, Decision Trees is a good choice since there are lots of independent variables, and this model can handle all of them easily to generate efficient rules.

3- AdaBoost

- this model could apply to determine students pass or fail in a course
- Boost model is working sequentially and every time it learns from the previous trained

- one of the disadvantages of AdaBoost is a weak learner

Overall, AdaBoost is a good choice even if it requires a longer time to train the data since it can deal with outliers and fit all the points perfectly.

''' \*\*

## 0.6.5 Implementation - Creating a Training and Predicting Pipeline

To properly evaluate the performance of each model you've chosen, it's important that you create a training and predicting pipeline that allows you to quickly and effectively train models using various sizes of training data and perform predictions on the testing data. Your implementation here will be used in the following section. In the code block below, you will need to implement the following: - Import `fbeta_score` and `accuracy_score` from `sklearn.metrics`. - Fit the learner to the sampled training data and record the training time. - Perform predictions on the test data `X_test`, and also on the first 300 training points `X_train[:300]`. - Record the total prediction time. - Calculate the accuracy score for both the training subset and testing set. - Calculate the F-score for both the training subset and testing set. - Make sure that you set the beta parameter!

```
In [37]: # TODO: Import two metrics from sklearn - fbeta_score and accuracy_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import fbeta_score
def train_predict(learner, sample_size, X_train, y_train, X_test, y_test):
    '''
    inputs:
        - learner: the learning algorithm to be trained and predicted on
        - sample_size: the size of samples (number) to be drawn from training set
        - X_train: features training set
        - y_train: income training set
        - X_test: features testing set
        - y_test: income testing set
    '''

    results = {}

    # TODO: Fit the learner to the training data using slicing with 'sample_size' using
    start = time() # Get start time
    learner = learner.fit(X_train[: sample_size], y_train[: sample_size])
    end = time() # Get end time

    # TODO: Calculate the training time
    results['train_time'] = end - start

    # TODO: Get the predictions on the test set(X_test),
    #         then get predictions on the first 300 training samples(X_train) using .pred
    start = time() # Get start time
    predictions_test = learner.predict(X_test)
    predictions_train = learner.predict(X_train[:300])
    end = time() # Get end time
```

```

# TODO: Calculate the total prediction time
results['pred_time'] = end - start

# TODO: Compute accuracy on the first 300 training samples which is y_train[:300]
results['acc_train'] = accuracy_score(y_train[:300], predictions_train)

# TODO: Compute accuracy on test set using accuracy_score()
results['acc_test'] = accuracy_score(y_test, predictions_test)

# TODO: Compute F-score on the the first 300 training samples using fbeta_score()
results['f_train'] = fbeta_score(y_train[:300], predictions_train, beta=0.5)

# TODO: Compute F-score on the test set which is y_test
results['f_test'] = fbeta_score(y_test, predictions_test, beta=0.5)

# Success
print("{} trained on {} samples.".format(learner.__class__.__name__, sample_size))

# Return the results
return results

```

## 0.6.6 Implementation: Initial Model Evaluation

In the code cell, you will need to implement the following: - Import the three supervised learning models you've discussed in the previous section. - Initialize the three models and store them in 'clf\_A', 'clf\_B', and 'clf\_C'. - Use a 'random\_state' for each model you use, if provided. - **Note:** Use the default settings for each model — you will tune one specific model in a later section. - Calculate the number of records equal to 1%, 10%, and 100% of the training data. - Store those values in 'samples\_1', 'samples\_10', and 'samples\_100' respectively.

**Note:** Depending on which algorithms you chose, the following implementation may take some time to run!

```

In [40]: # TODO: Import the three supervised learning models from sklearn
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
# TODO: Initialize the three models
clf_A = LogisticRegression(random_state=42)
clf_B = DecisionTreeClassifier(random_state=42)
clf_C = AdaBoostClassifier(random_state=42)

# TODO: Calculate the number of samples for 1%, 10%, and 100% of the training data
# HINT: samples_100 is the entire training set i.e. len(y_train)
# HINT: samples_10 is 10% of samples_100 (ensure to set the count of the values to be `
# HINT: samples_1 is 1% of samples_100 (ensure to set the count of the values to be `in
samples_100 = len(y_train)
samples_10 = int(len(y_train)*(10/100))

```

```

samples_1 = int(len(y_train)*(1/100))

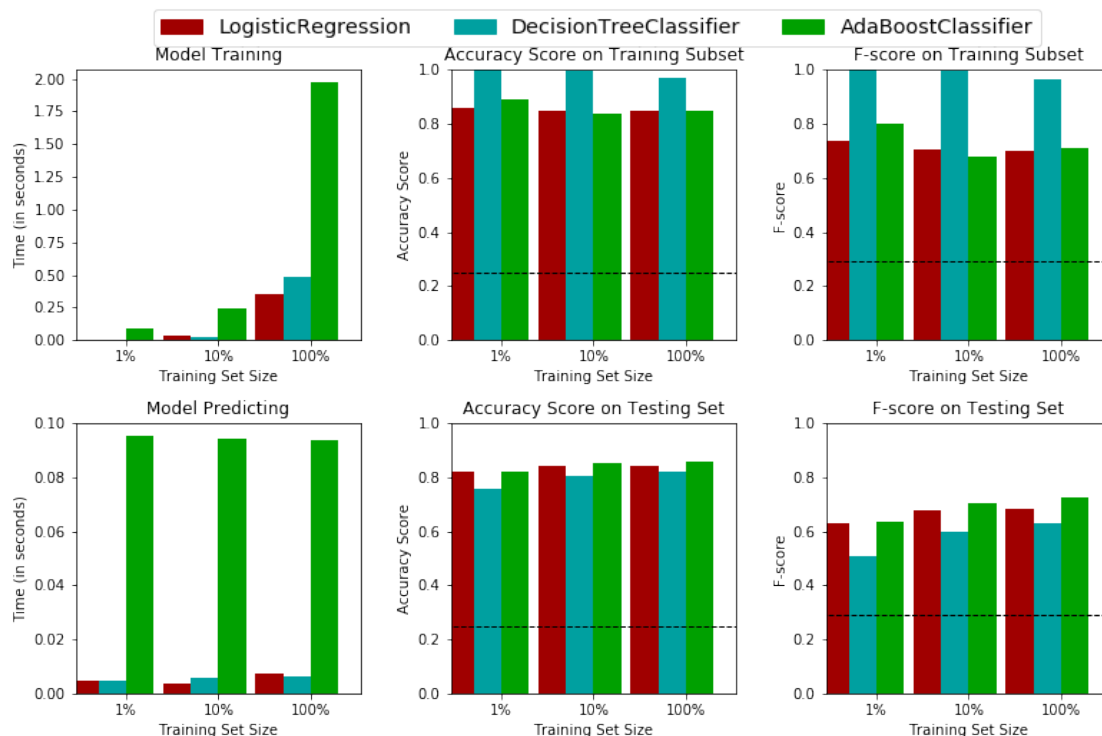
# Collect results on the learners
results = {}
for clf in [clf_A, clf_B, clf_C]:
    clf_name = clf.__class__.__name__
    results[clf_name] = {}
    for i, samples in enumerate([samples_1, samples_10, samples_100]):
        results[clf_name][i] = \
            train_predict(clf, samples, X_train, y_train, X_test, y_test)

# Run metrics visualization for the three supervised learning models chosen
vs.evaluate(results, accuracy, fscore)

```

LogisticRegression trained on 361 samples.  
 LogisticRegression trained on 3617 samples.  
 LogisticRegression trained on 36177 samples.  
 DecisionTreeClassifier trained on 361 samples.  
 DecisionTreeClassifier trained on 3617 samples.  
 DecisionTreeClassifier trained on 36177 samples.  
 AdaBoostClassifier trained on 361 samples.  
 AdaBoostClassifier trained on 3617 samples.  
 AdaBoostClassifier trained on 36177 samples.

Performance Metrics for Three Supervised Learning Models



---

## 0.7 Improving Results

In this final section, you will choose from the three supervised learning models the *best* model to use on the student data. You will then perform a grid search optimization for the model over the entire training set (`X_train` and `y_train`) by tuning at least one parameter to improve upon the untuned model's F-score.

### 0.7.1 Question 3 - Choosing the Best Model

- Based on the evaluation you performed earlier, in one to two paragraphs, explain to *CharityML* which of the three models you believe to be most appropriate for the task of identifying individuals that make more than \$50,000.

**HINT:** Look at the graph at the bottom left from the cell above (the visualization created by `vs.evaluate(results, accuracy, fscore)`) and check the F score for the testing set when 100% of the training set is used. Which model has the highest score? Your answer should include discussion of the: \* metrics - F score on the testing when 100% of the training data is used, \* prediction/training time \* the algorithm's suitability for the data.

**Answer:**

- the highest F-score when 100% of the training data is used is for AdaBoost classifier
- logistic regression is the fastest model then Decision tree respectively
- the most appropriate model is the Adaboost classifier because it shows the highest F-score in testing data

**\*\***

### 0.7.2 Question 4 - Describing the Model in Layman's Terms

- In one to two paragraphs, explain to *CharityML*, in layman's terms, how the final model chosen is supposed to work. Be sure that you are describing the major qualities of the model, such as how the model is trained and how the model makes a prediction. Avoid using advanced mathematical jargon, such as describing equations.

**HINT:**

When explaining your model, if using external resources please include all citations. **Answer:** The AdaBoost classifier is a strong classifier that builds from multiple weak classifiers. It works stochastically, and each time the new model corrects the previous model until it becomes strong and efficient.

AdaBoost is a very successful algorithm developed for binary classification so, it will support the work.

It works as the following: 1- initialize equal weights for all the data point 2- increase the weight for the misclassified point after deploying the first model 3- repeat the second step and each time increase the weight for misclassified points 4- contain all the models together to become one strong model

Reference: Boosting in Machine Learning, Geeksforgeeks.org, 03-May-2019. [Online]. Available: <https://www.geeksforgeeks.org/boosting-in-machine-learning-boosting-and-adaboost/>. [Accessed: 12-Jun-2021].\*\*

### 0.7.3 Implementation: Model Tuning

Fine tune the chosen model. Use grid search (GridSearchCV) with at least one important parameter tuned with at least 3 different values. You will need to use the entire training set for this. In the code cell below, you will need to implement the following: - Import `sklearn.grid_search.GridSearchCV` and `sklearn.metrics.make_scorer`. - Initialize the classifier you've chosen and store it in `clf`. - Set a `random_state` if one is available to the same state you set before. - Create a dictionary of parameters you wish to tune for the chosen model. - Example: `parameters = {'parameter' : [list of values]}`. - **Note:** Avoid tuning the `max_features` parameter of your learner if that parameter is available! - Use `make_scorer` to create an `fbeta_score` scoring object (with  $\beta = 0.5$ ). - Perform grid search on the classifier `clf` using the 'scorer', and store it in `grid_obj`. - Fit the grid search object to the training data (`X_train`, `y_train`), and store it in `grid_fit`.

**Note:** Depending on the algorithm chosen and the parameter list, the following implementation may take some time to run!

```
In [41]: # TODO: Import 'GridSearchCV', 'make_scorer', and any other necessary libraries
from sklearn.metrics import make_scorer
from sklearn.model_selection import GridSearchCV
# TODO: Initialize the classifier
clf = clf_C

# TODO: Create the parameters list you wish to tune, using a dictionary if needed.
# HINT: parameters = {'parameter_1': [value1, value2], 'parameter_2': [value1, value2]}
parameters = {'n_estimators': [25, 50, 75, 100], 'learning_rate': [0.1, 1, 1.5, 2], 'algorithm': [

# TODO: Make an fbeta_score scoring object using make_scorer()
scorer = make_scorer(fbeta_score, beta=0.5)

# TODO: Perform grid search on the classifier using 'scorer' as the scoring method using
grid_obj = GridSearchCV(clf, parameters, scoring=scorer)

# TODO: Fit the grid search object to the training data and find the optimal parameters
grid_fit = grid_obj.fit(X_train, y_train)

# Get the estimator
best_clf = grid_fit.best_estimator_

# Make predictions using the unoptimized and model
predictions = (clf.fit(X_train, y_train)).predict(X_test)
best_predictions = best_clf.predict(X_test)

# Report the before-and-afterscores
print("Unoptimized model\n-----")
```

```

print("Accuracy score on testing data: {:.4f}".format(accuracy_score(y_test, prediction
print("F-score on testing data: {:.4f}".format(fbeta_score(y_test, predictions, beta =
print("\nOptimized Model\n-----")
print("Final accuracy score on the testing data: {:.4f}".format(accuracy_score(y_test,
print("Final F-score on the testing data: {:.4f}".format(fbeta_score(y_test, best_predi

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWa
'precision', 'predicted', average, warn_for)

```

Unoptimized model

-----

Accuracy score on testing data: 0.8576

F-score on testing data: 0.7246

Optimized Model

-----

Final accuracy score on the testing data: 0.8647

Final F-score on the testing data: 0.7382

#### 0.7.4 Question 5 - Final Model Evaluation

- What is your optimized model's accuracy and F-score on the testing data?
- Are these scores better or worse than the unoptimized model?
- How do the results from your optimized model compare to the naive predictor benchmarks you found earlier in **Question 1**?\_

**Note:** Fill in the table below with your results, and then provide discussion in the **Answer** box.

Metric	Unoptimized Model	Optimized Model
Accuracy Score	0.8576	0.8647
F-score	0.7246	0.7382

**Results:** Answer: the accuracy score in Optimized model is greater than Unoptimized Model also the F-score is improved in Optimized Model

### 0.8 Feature Importance

An important task when performing supervised learning on a dataset like the census data we study here is determining which features provide the most predictive power. By focusing on the relationship between only a few crucial features and the target label we simplify our understanding of the phenomenon, which is most always a useful thing to do. In the case of this project, that means we wish to identify a small number of features that most strongly predict whether an individual makes at most or more than \$50,000.

Choose a scikit-learn classifier (e.g., adaboost, random forests) that has a `feature_importance_` attribute, which is a function that ranks the importance of features according to the chosen classifier. In the next python cell fit this classifier to training set and use this attribute to determine the top 5 most important features for the census dataset.

#### 0.8.1 Question 6 - Feature Relevance Observation

When **Exploring the Data**, it was shown there are thirteen available features for each individual on record in the census data. Of these thirteen records, which five features do you believe to be most important for prediction, and in what order would you rank them and why? the most important features in my opinion in order are:

1- Work class 2- Native country 3- Hours per week 4- Occupation 5- Sex

#### 0.8.2 Implementation - Extracting Feature Importance

Choose a scikit-learn supervised learning algorithm that has a `feature_importance_` attribute available for it. This attribute is a function that ranks the importance of each feature when making predictions based on the chosen algorithm.

In the code cell below, you will need to implement the following: - Import a supervised learning model from sklearn if it is different from the three used earlier. - Train the supervised model on the entire training set. - Extract the feature importances using `'feature_importances_'`.




```
In [42]: # TODO: Import a supervised learning model that has 'feature_importances_'
from sklearn.ensemble import RandomForestClassifier

# TODO: Train the supervised model on the training set using .fit(X_train, y_train)
model = RandomForestClassifier()
model.fit(X_train, y_train)
p= model.predict(X_test)
print(accuracy_score(y_test, p))
# TODO: Extract the feature importances using .feature_importances_
importances = model.feature_importances_

# Plot
vs.feature_plot(importances, X_train, y_train)
```

0.839137645108



### 0.8.3 Question 7 - Extracting Feature Importance

Observe the visualization created above which displays the five most relevant features for predicting if an individual makes at most or above \$50,000.

\* How do these five features compare to the five features you discussed in **Question 6**? \* If you were close to the same answer, how does this visualization confirm your thoughts? \* If you were not close, why do you think these features are more relevant? I choose work class, occupation, and hours of a week since they determine how much the worker will gain. And they are different from one country to another and from gender to another

The figure shows different features than I was thought. These features are more relevant because they have the most weights

### 0.8.4 Feature Selection

How does a model perform if we only use a subset of all the available features in the data? With less features required to train, the expectation is that training and prediction time is much lower — at the cost of performance metrics. From the visualization above, we see that the top five most important features contribute more than half of the importance of **all** features present in the data. This hints that we can attempt to *reduce the feature space* and simplify the information required for the model to learn. The code cell below will use the same optimized model you found earlier, and train it on the same training set *with only the top five important features*.

```
In [43]: # Import functionality for cloning a model
        from sklearn.base import clone

        # Reduce the feature space
        X_train_reduced = X_train[X_train.columns.values[(np.argsort(importances)[:,-1])[:5]]]
        X_test_reduced = X_test[X_test.columns.values[(np.argsort(importances)[:,-1])[:5]]]

        # Train on the "best" model found from grid search earlier
        clf = (clone(best_clf)).fit(X_train_reduced, y_train)

        # Make new predictions
        reduced_predictions = clf.predict(X_test_reduced)

        # Report scores from the final model using both versions of data
        print("Final Model trained on full data\n-----")
        print("Accuracy on testing data: {:.4f}".format(accuracy_score(y_test, best_predictions)))
        print("F-score on testing data: {:.4f}".format(fbeta_score(y_test, best_predictions, beta=2)))
        print("\nFinal Model trained on reduced data\n-----")
        print("Accuracy on testing data: {:.4f}".format(accuracy_score(y_test, reduced_predictions)))
        print("F-score on testing data: {:.4f}".format(fbeta_score(y_test, reduced_predictions, beta=2)))
```

Final Model trained on full data

-----

Accuracy on testing data: 0.8647

F-score on testing data: 0.7382

Final Model trained on reduced data

-----

Accuracy on testing data: 0.8495

F-score on testing data: 0.7067

### 0.8.5 Question 8 - Effects of Feature Selection

- How does the final model's F-score and accuracy score on the reduced data using only five features compare to those same scores when all features are used?
- If training time was a factor, would you consider using the reduced data as your training set?

the Accuracy score and F-score are less in the reduced data model, and if the time was a factor, I prefer to choose the reduced data model

**Note:** Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to

**File -> Download as -> HTML (.html).** Include the finished document along with this notebook as your submission.

## 0.9 Before You Submit

You will also need run the following in order to convert the Jupyter notebook into HTML, so that your submission will include both files.

```
In [44]: !!jupyter nbconvert *.ipynb
```

```
Out[44]: ['[NbConvertApp] Converting notebook finding_donors.ipynb to html',  
          '[NbConvertApp] Writing 532290 bytes to finding_donors.html']
```